

Learner Notebook

July 7, 2025

1 Importing Libraries

```
[7]: %matplotlib inline

import os
import shutil
import random
import torch
import torchvision
import numpy as np

from PIL import Image
from matplotlib import pyplot as plt

torch.manual_seed(0)

print('Using PyTorch version', torch.__version__)
```

Using PyTorch version 2.3.1+cu121

2 Preparing Training and Test Sets

```
[8]: class_names = ['normal', 'viral', 'covid']
root_dir = 'COVID-19 Radiography Database'
source_dirs = ['NORMAL', 'Viral Pneumonia', 'COVID-19']

if os.path.isdir(os.path.join(root_dir, source_dirs[1])):
    os.mkdir(os.path.join(root_dir, 'test'))

    for i, d in enumerate(source_dirs):
        os.rename(os.path.join(root_dir, d), os.path.join(root_dir,
↵class_names[i]))

    for c in class_names:
        os.mkdir(os.path.join(root_dir, 'test', c))

    for c in class_names:
```

```

        images = [x for x in os.listdir(os.path.join(root_dir, c)) if x.lower().
↪endswith('png')]
        selected_images = random.sample(images, 30)
        for image in selected_images:
            source_path = os.path.join(root_dir, c, image)
            target_path = os.path.join(root_dir, 'test', c, image)
            shutil.move(source_path, target_path)

```

3 Creating Custom Dataset

```

[9]: class ChestXRayDataset(torch.utils.data.Dataset):
    def __init__(self, image_dirs, transform):
        def get_images(class_name):
            images = [x for x in os.listdir(image_dirs[class_name]) if x.
↪lower().endswith('png')]
            print(f'Found {len(images)} {class_name}')
            return images
        self.images={}
        self.class_names=['normal', 'viral', 'covid']
        for c in self.class_names:
            self.images[c]=get_images(c)
        self.image_dirs=image_dirs
        self.transform=transform
    def __len__(self):
        return sum([len(self.images[c]) for c in self.class_names])
    def __getitem__(self, index):
        class_name=random.choice(self.class_names)
        index=index%len(self.images[class_name])
        image_name=self.images[class_name][index]
        image_path=os.path.join(self.image_dirs[class_name], image_name)
        image=Image.open(image_path).convert('RGB')
        return self.transform(image), self.class_names.index(class_name)

```

4 Image Transformations

```

[10]: train_transform = torchvision.transforms.Compose([
    torchvision.transforms.Resize(size=(224,224)),
    torchvision.transforms.RandomHorizontalFlip(),
    torchvision.transforms.ToTensor(),
    torchvision.transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229,0.
↪224,0.225])
])
test_transform = torchvision.transforms.Compose([
    torchvision.transforms.Resize(size=(224,224)),
    torchvision.transforms.ToTensor(),

```

```

        torchvision.transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229,0.
↪224,0.225])
    ])

```

5 Prepare DataLoader

```

[11]: train_dirs = {
        'normal': 'COVID-19 Radiography Database/normal',
        'viral': 'COVID-19 Radiography Database/viral',
        'covid': 'COVID-19 Radiography Database/covid'
    }
    train_dataset=ChestXRayDataset(train_dirs, train_transform)

```

```

Found 1311normal
Found 1315viral
Found 189covid

```

```

[12]: test_dirs = {
        'normal': 'COVID-19 Radiography Database/test/normal',
        'viral': 'COVID-19 Radiography Database/test/viral',
        'covid': 'COVID-19 Radiography Database/test/covid'
    }
    test_dataset = ChestXRayDataset(test_dirs, test_transform)

```

```

Found 30normal
Found 30viral
Found 30covid

```

```

[13]: batch_size=6
    dl_train = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size,
↪shuffle=True)
    dl_test = torch.utils.data.DataLoader(test_dataset, batch_size=batch_size,
↪shuffle=True)
    print('Num of training batches', len(dl_train))
    print('Num of test batches', len(dl_test))

```

```

Num of training batches 470
Num of test batches 15

```

6 Data Visualization

```

[14]: class_names=train_dataset.class_names
    def show_images(images, labels, preds):
        plt.figure(figsize=(8,4))
        for i, image in enumerate(images):
            plt.subplot(1,6,i+1, xticks=[], yticks=[])

```

```


image=image.numpy().transpose((1,2,0))
mean=np.array([0.485,0.456,0.406])
std= np.array([0.229, 0.224, 0.225])
image=image*std/mean
image=np.clip(image,0.,1.)
plt.imshow(image)
col = 'green' if preds[i]==labels[i] else 'red'
plt.xlabel(f'{class_names[int(labels[i].numpy())]}')
plt.ylabel(f'{class_names[int(preds[i].numpy())]}', color=col)
plt.tight_layout()
plt.show()

```

```

[15]: images, labels =next(iter(dl_train))
show_images(images, labels, labels)


```



```

[16]: images, labels =next(iter(dl_test))
show_images(images, labels, labels)

```



7 Creating the Model

```

[17]: resnet18 =torchvision.models.resnet18(pretrained=True)
print(resnet18)

```

```

ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
    bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
    track_running_stats=True)
  (relu): ReLU(inplace=True)

```

```

(maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
ceil_mode=False)
(layer1): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
  (1): BasicBlock(
    (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
(layer2): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

        (relu): ReLU(inplace=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(layer3): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
(layer4): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (downsample): Sequential(

```

```

        (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(1): BasicBlock(
  (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
  (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (relu): ReLU(inplace=True)
  (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
  (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
)
)
(avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
(fc): Linear(in_features=512, out_features=1000, bias=True)
)

```

```

[18]: resnet18.fc=torch.nn.Linear(in_features=512, out_features=3)
      loss_fn=torch.nn.CrossEntropyLoss()
      optimizer=torch.optim.Adam(resnet18.parameters(), lr=3e-5)

```

```


[19]: def show_preds():
      resnet18.eval()
      images, labels =next(iter(dl_test))
      outputs = resnet18(images)
      _, preds=torch.max(outputs, 1)
      show_images(images, labels, preds)

```

```

[20]: show_preds()

```




8 Training the Model


```
[27]: def train(epochs):
    print('Starting training..')
    for e in range(0, epochs):
        print(f'Starting epoch {e+1}/{epochs}')
        print('='*20)
        train_loss=0
        resnet18.train()
        for train_step, (images, labels) in enumerate(dl_train):
            optimizer.zero_grad()
            outputs=resnet18(images)
            loss=loss_fn(outputs, labels)
            loss.backward()
            optimizer.step()
            train_loss=loss.item()
            if train_step%20==0:
                print('Evaluating at step', train_step)
                acc=0.
                val_loss=0.
                resnet18.eval()
                for val_step, (images, labels) in enumerate(dl_test):
                    outputs=resnet18(images)
                    loss=loss_fn(outputs, labels)
                    val_loss+=loss.item()
                    _,preds=torch.max(outputs, 1)
                    acc+=sum(preds==labels).numpy()
                val_loss/=(val_step+1)
                acc=acc/len(test_dataset)
                print(f'Val loss: {val_loss:.4f}, Acc: {acc:.4f}')
                show_preds()
                resnet18.train()
                if acc > 0.95:
                    print('Performance condition satisfied..')
                    return
            train_loss/=(train_step+1)
        print(f'Training loss: {train_loss:.4f}')
```

```
[28]: train(epochs=1)
```

```
Starting training..
Starting epoch 1/1
=====
Evaluating at step 0
Val loss: 0.2160, Acc: 0.9444
```

Evaluating at step 20
Val loss: 0.1231, Acc: 0.9556




Performance condition satisfied..


[28]: `train(epochs=1)`

Starting training..
Starting epoch 1/1
=====

Evaluating at step 0
Val loss: 0.2160, Acc: 0.9444



Evaluating at step 20
Val loss: 0.1231, Acc: 0.9556



Performance condition satisfied..

9 Final Results

[26]: `show_preds()`

